

## Model Optimization and Tuning Report

Date	15 April 2024
Team ID	Team-738164
Project Title	Rainfall Prediction Using Machine Learning
Maximum Marks	10 Marks

### Model Optimization and Tuning Report:

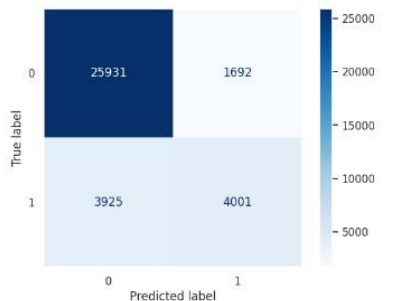
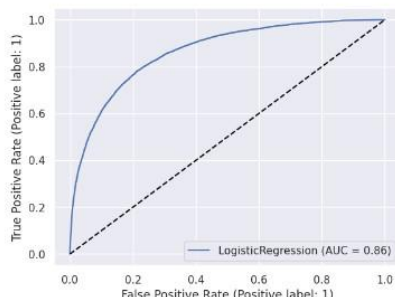
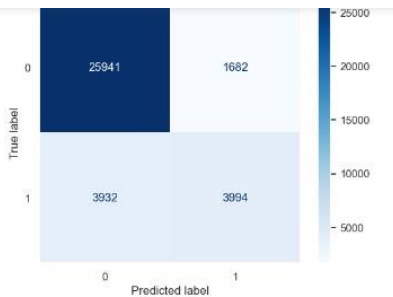
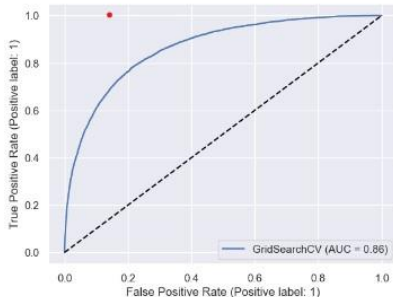
The Model Optimization and Tuning Report involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

### Hyperparameter Tuning Documentation (6 Marks):


Model	Tuned Hyperparameters	Optimal Values
Logistic Regression	<p><b>Hyperparameter Tuning</b></p> <pre>In [84]: logreg_params = {     'C': [1, 1e8, 1e16],     'fit_intercept': [True, False],     'max_iter': [50, 100, 150],     'random_state': [42] }  logreg_gs = GridSearchCV(logreg, logreg_params, scoring='accuracy', n_jobs=-1, cv=3) logreg_gs.fit(X_train, y_train)</pre>	<pre>In [89]: logreg_gs.best_params_ Out[89]: {'C': 10000000.0, 'fit_intercept': True, 'max_iter': 50, 'random_state': 42}</pre> <p>checking model fitness</p> <pre>----- Train score: 0.8469 Test score: 0.8421</pre>
Random Forest	<p><b>Hyperparameter Tuning</b></p> <pre>In [104]: rf_params = {     'n_estimators': [10, 35, 100],     'criterion': ['gini', 'entropy'],     'max_depth': [3, 7, 11],     'min_samples_split': [2, 5, 10],     'min_samples_leaf': [1, 3, 5],     'random_state': [42] }  rf_gs = GridSearchCV(rf, param_grid=rf_params, scoring='accuracy', n_jobs=-1, cv=3) rf_gs.fit(X_train, y_train)</pre>	<pre>In [107]: rf_gs.best_params_ Out[107]: {'criterion': 'gini',     'max_depth': 11,     'min_samples_leaf': 1,     'min_samples_split': 2,     'n_estimators': 100,     'random_state': 42}</pre> <p>checking model fitness</p> <pre>----- Train score: 0.8821 Test score: 0.8493</pre>

<h2>Decision Tree</h2>	<p><b>Hyperparameter Tuning</b></p> <pre>In [95]: params = {         'criterion': ['gini', 'entropy'],         'max_depth': [3, 7, 11],         'min_samples_split': [2, 5, 10],         'min_samples_leaf': [3, 5, 5],         'random_state': [42]     }      clf_gs = GridSearchCV(clf, param_grid=params, scoring='accuracy', n_jobs=-1, cv=3)     clf_gs.fit(X_train, y_train)</pre>	<pre>In [98]: clf_gs.best_params_  Out[98]: {'criterion': 'gini',          'max_depth': 7,          'min_samples_leaf': 5,          'min_samples_split': 2,          'random_state': 42}</pre> <p>Checking model fitness</p> <hr/> <p>Train score: 0.8475 Test score: 0.8405</p>
<h2>XG-Boost</h2>	<p><b>Hyperparameter Tuning</b></p> <pre>In [113]: xgb_params = {         'n_estimators': [10, 35, 100],         'max_depth': [5, 10, 15],         'learning_rate': [0.01, 0.1, 0.25]     }      xgb_gs = GridSearchCV(xgb, xgb_params, scoring='accuracy', n_jobs=-1, cv=3)     xgb_gs.fit(X_train, y_train)  Out[113]: GridSearchCV(cv=3,                     estimator=XGBClassifier(base_score=None, booster=None,                     callbacks=None, colsample_bylevel=None,                     colsample_bynode=None, device=None,                     colsample_bytree=None, early_stopping_rounds=None,                     enable_categorical=False, eval_metric=None,                     feature_types=None, gamma=None,                     grow_policy=None, importance_type=None,                     interaction_constraints=None,                     learning_rate=None,                     max_cat_to_onehot=None,                     max_delta_step=None, max_depth=None,                     max_leaves=None, min_child_weight=None,                     missing=None, monotone_constraints=None,                     multi_strategy=None, n_estimators=None,                     n_jobs=None, num_parallel_tree=None,                     random_state=None,                     n_jobs=-1,                     param_grid={'learning_rate': [0.01, 0.1, 0.25],                     'max_depth': [5, 10, 15],                     'n_estimators': [10, 35, 100]},                     scoring='accuracy')</pre>	<pre>In [116]: xgb_gs.best_params_  Out[116]: {'learning_rate': 0.1, 'max_depth': 10, 'n_estimators': 100}</pre> <p>Checking model fitness</p> <hr/> <p>Train score: 0.9329 Test score: 0.8616</p>

## Performance Metrics Comparison Report (2 Marks):


Model	Baseline Metric	Optimized Metric																																																																														
Logistic Regression	<div><p>Confusion Matrix for Baseline Logistic Regression:</p><table><tr><th>True label \ Predicted label</th><th>0</th><th>1</th></tr><tr><th>0</th><td>25931</td><td>1692</td></tr><tr><th>1</th><td>3925</td><td>4001</td></tr></table><p>Classification Report</p><table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>0.87</td><td>0.94</td><td>0.90</td><td>27623</td></tr><tr><td>1</td><td>0.70</td><td>0.50</td><td>0.59</td><td>7926</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.84</td><td>35549</td></tr><tr><td>macro avg</td><td>0.79</td><td>0.72</td><td>0.74</td><td>35549</td></tr><tr><td>weighted avg</td><td>0.83</td><td>0.84</td><td>0.83</td><td>35549</td></tr></table><p>ROC Curve</p><p>ROC Curve for Baseline Logistic Regression: AUC = 0.86</p><p>Checking model fitness</p><p>Train score: 0.8468 Test score: 0.842</p></div> <div>Accuracy: 84%    AUC: 0.86</div>	True label \ Predicted label	0	1	0	25931	1692	1	3925	4001		precision	recall	f1-score	support	0	0.87	0.94	0.90	27623	1	0.70	0.50	0.59	7926	accuracy			0.84	35549	macro avg	0.79	0.72	0.74	35549	weighted avg	0.83	0.84	0.83	35549	<div><p>Confusion Matrix for Optimized Logistic Regression:</p><table><tr><th>True label \ Predicted label</th><th>0</th><th>1</th></tr><tr><th>0</th><td>25941</td><td>1682</td></tr><tr><th>1</th><td>3932</td><td>3994</td></tr></table><p>Classification Report</p><table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>0.87</td><td>0.94</td><td>0.90</td><td>27623</td></tr><tr><td>1</td><td>0.70</td><td>0.50</td><td>0.59</td><td>7926</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.84</td><td>35549</td></tr><tr><td>macro avg</td><td>0.79</td><td>0.72</td><td>0.74</td><td>35549</td></tr><tr><td>weighted avg</td><td>0.83</td><td>0.84</td><td>0.83</td><td>35549</td></tr></table><p>ROC Curve</p><p>ROC Curve for Optimized Logistic Regression: AUC = 0.86</p><p>Checking model fitness</p><p>Train score: 0.8469 Test score: 0.8421</p></div> <div>Accuracy: 84%    AUC: 0.86</div>	True label \ Predicted label	0	1	0	25941	1682	1	3932	3994		precision	recall	f1-score	support	0	0.87	0.94	0.90	27623	1	0.70	0.50	0.59	7926	accuracy			0.84	35549	macro avg	0.79	0.72	0.74	35549	weighted avg	0.83	0.84	0.83	35549
True label \ Predicted label	0	1																																																																														
0	25931	1692																																																																														
1	3925	4001																																																																														
	precision	recall	f1-score	support																																																																												
0	0.87	0.94	0.90	27623																																																																												
1	0.70	0.50	0.59	7926																																																																												
accuracy			0.84	35549																																																																												
macro avg	0.79	0.72	0.74	35549																																																																												
weighted avg	0.83	0.84	0.83	35549																																																																												
True label \ Predicted label	0	1																																																																														
0	25941	1682																																																																														
1	3932	3994																																																																														
	precision	recall	f1-score	support																																																																												
0	0.87	0.94	0.90	27623																																																																												
1	0.70	0.50	0.59	7926																																																																												
accuracy			0.84	35549																																																																												
macro avg	0.79	0.72	0.74	35549																																																																												
weighted avg	0.83	0.84	0.83	35549																																																																												

## Random Forest



Classification Report


	precision	recall	f1-score	support
0	0.87	0.95	0.91	27623
1	0.75	0.51	0.61	7926
accuracy				0.85 35549
macro avg	0.81	0.73	0.76	35549
weighted avg	0.85	0.85	0.84	35549



Checking model fitness

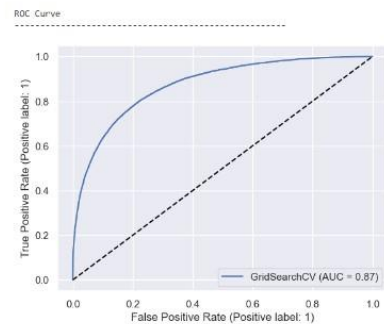
Train score:	1.0
Test score:	0.8538

Accuracy: 85%    AUC: 0.88



Classification Report

	precision	recall	f1-score	support
0	0.87	0.95	0.91	27623
1	0.75	0.48	0.59	7926
accuracy				0.85 35549
macro avg	0.81	0.72	0.75	35549
weighted avg	0.84	0.85	0.84	35549

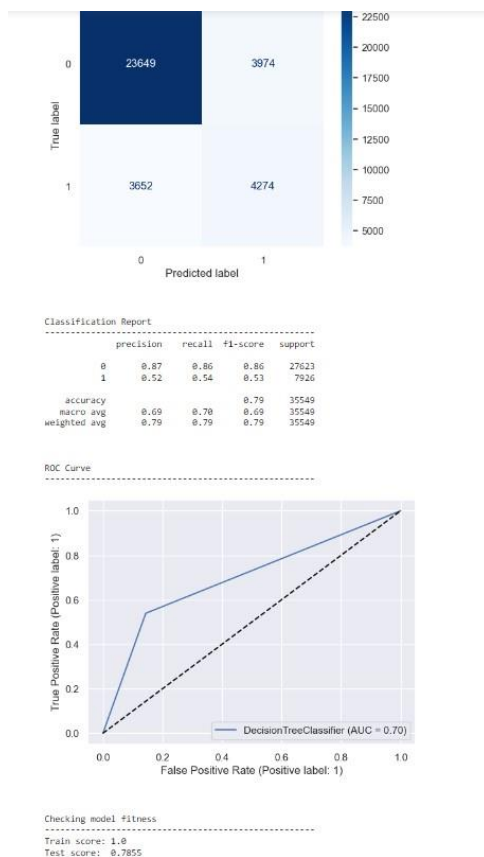


Checking model fitness

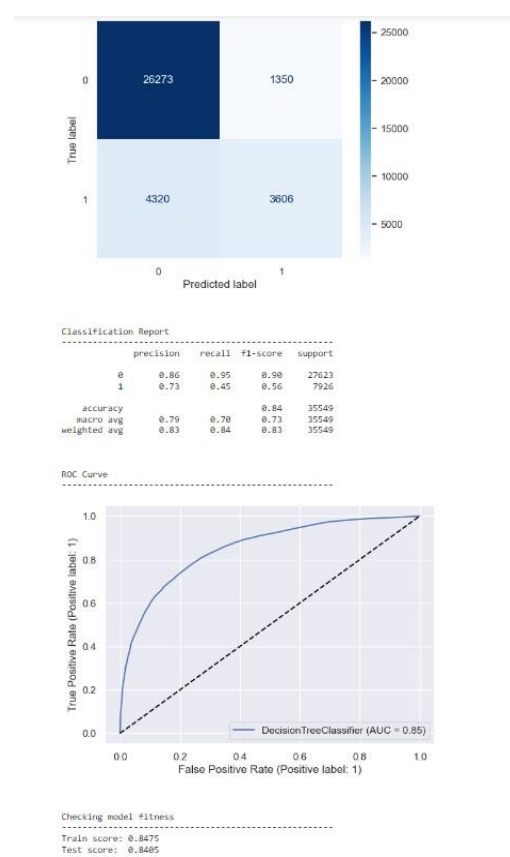
Train score:	0.8821
Test score:	0.8493

Accuracy: 85%    AUC: 0.88

## Decision Tree

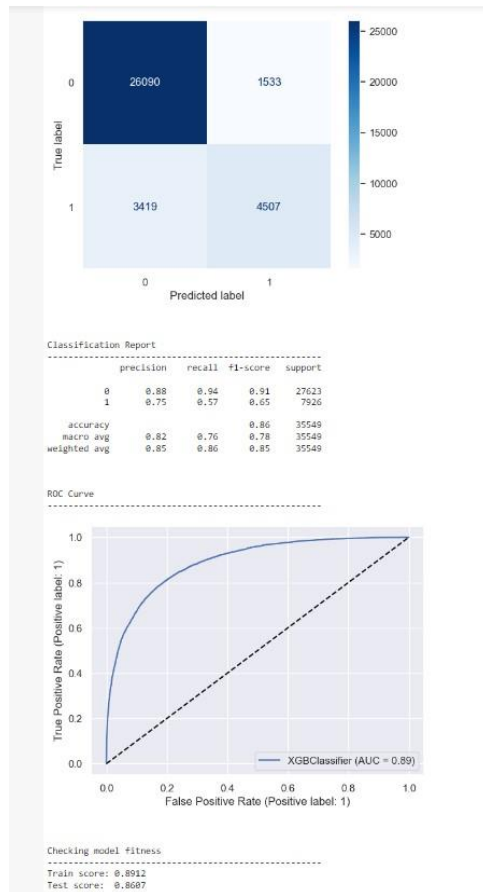


Accuracy: 79%    AUC: 0.70

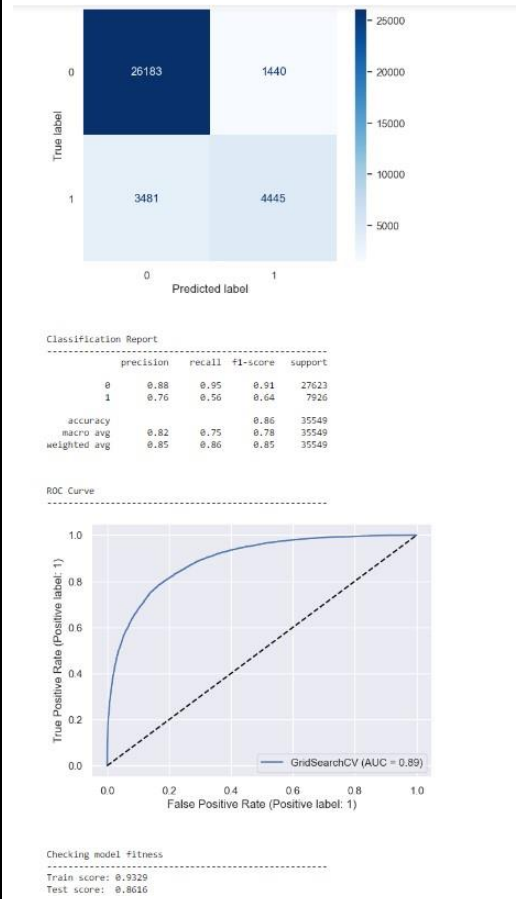


Accuracy: 84%    AUC: 0.85

## XG-Boost



Accuracy: 86%    AUC: 0.89



Accuracy: 86%    AUC: 0.89

**Final Model Selection Justification (2 Marks):**

Final Model	Reasoning
XG-Boost	The best performing model is the hyperparameter-tuned XG-Boost model with an accuracy of approximately 86%. The scores for both the training and testing data were similar, reducing concerns of the model being overfit.