

## Model Optimization and Tuning Report

Date	15 April 2024
Team ID	Team-738164
Project Title	Rainfall Prediction Using Machine Learning
Maximum Marks	10 Marks

### Model Optimization and Tuning Phase:

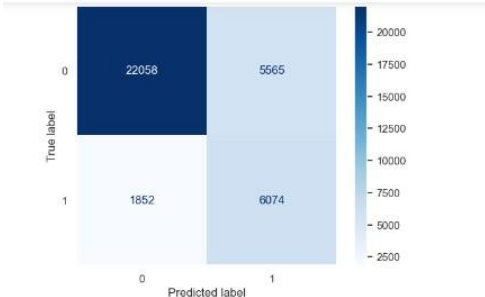
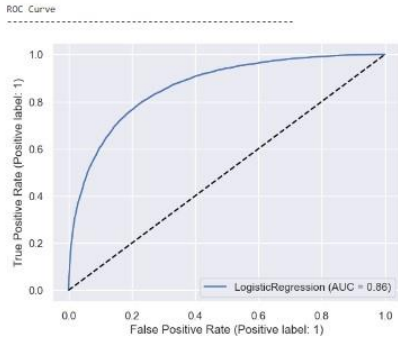

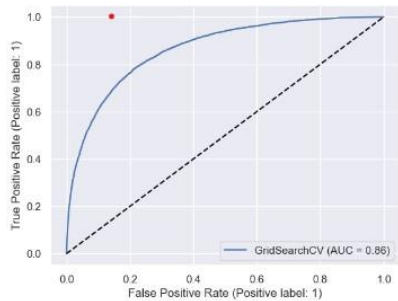
The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

### Hyperparameter Tuning Documentation (6 Marks):

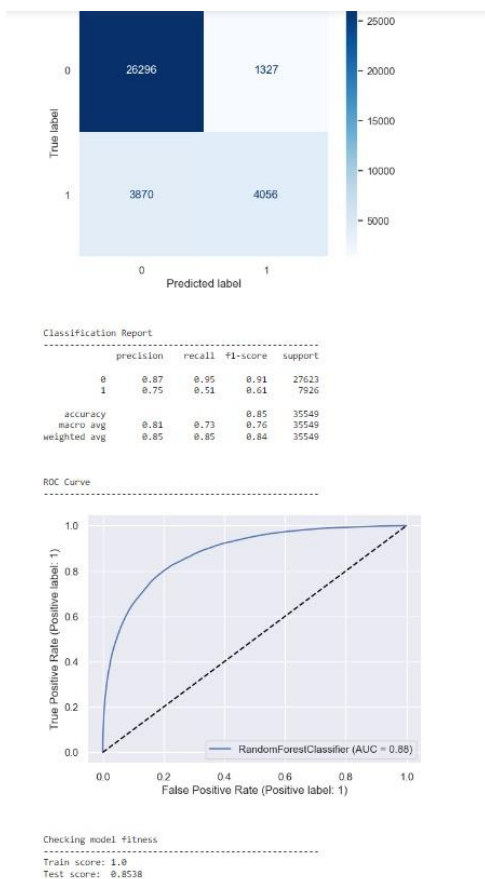
Model	Tuned Hyperparameters	Optimal Values
Logistic Regression	<p><b>Hyperparameter Tuning</b></p> <pre>In [84]: logreg_params = {     'C': [1, 1e8, 1e16],     'fit_intercept': [True, False],     'max_iter': [50, 100, 150],     'random_state': [42] }  logreg_gs = GridSearchCV(logreg, logreg_params, scoring='accuracy', n_jobs=-1, cv=3) logreg_gs.fit(X_train, y_train)</pre>	<pre>In [89]: logreg_gs.best_params_ Out[89]: {'C': 10000000.0, 'fit_intercept': True, 'max_iter': 50, 'random_state': 42}</pre> <p>Checking model fitness</p> <pre>----- Train score: 0.8469 Test score: 0.8421</pre>
Random Forest	<p><b>Hyperparameter Tuning</b></p> <pre>In [104]: rf_params = {     'n_estimators': [10, 35, 100],     'criterion': ['gini', 'entropy'],     'max_depth': [3, 7, 11],     'min_samples_split': [2, 5, 10],     'min_samples_leaf': [1, 3, 5],     'random_state': [42] }  rf_gs = GridSearchCV(rf, param_grid=rf_params, scoring='accuracy', n_jobs=-1, cv=3) rf_gs.fit(X_train, y_train)</pre>	<pre>In [107]: rf_gs.best_params_ Out[107]: {'criterion': 'gini',     'max_depth': 11,     'min_samples_leaf': 1,     'min_samples_split': 2,     'n_estimators': 100,     'random_state': 42}</pre> <p>Checking model fitness</p> <pre>----- Train score: 0.8821 Test score: 0.8493</pre>

<h2>Decision Tree</h2>	<p><b>Hyperparameter Tuning</b></p> <pre>In [95]: params = {         'criterion': ['gini', 'entropy'],         'max_depth': [3, 7, 11],         'min_samples_split': [4, 5, 10],         'min_samples_leaf': [3, 5, 9],         'random_state': [42]     }  clf_gs = GridSearchCV(clf, param_grid=params, scoring='accuracy', n_jobs=-1, cv=3) clf_gs.fit(X_train, y_train)</pre>	<pre>In [98]: clf_gs.best_params_  Out[98]: {'criterion': 'gini',         'max_depth': 7,         'min_samples_leaf': 5,         'min_samples_split': 2,         'random_state': 42}</pre> <p>Checking model fitness</p> <hr/> <p>Train score: 0.8475 Test score: 0.8405</p>
<h2>XG-Boost</h2>	<p><b>Hyperparameter Tuning</b></p> <pre>In [113]: xgb_params = {         'n_estimators': [10, 35, 100],         'max_depth': [5, 10, 15],         'learning_rate': [0.01, 0.1, 0.25]     }  xgb_gs = GridSearchCV(xgb, xgb_params, scoring='accuracy', n_jobs=-1, cv=3) xgb_gs.fit(X_train, y_train)</pre> <pre>Out[113]: GridSearchCV(cv=3,         estimator=XGBClassifier(base_score=None, booster=None,         callbacks=None, colsample_bylevel=None,         colsample_bynode=None, device=None,         colsample_bytree=None,         early_stopping_rounds=None,         enable_categorical=False, eval_metric=None,         feature_types=None, gamma=None,         grow_policy=None, importance_type=None,         interaction_constraints=None,         learning_rate=None,         max_cat_to_onehot=None,         max_delta_step=None, max_depth=None,         max_leaves=None, min_child_weight=None,         missing=None, monotone_constraints=None,         multi_strategy=None, n_estimators=None,         n_jobs=None, num_parallel_tree=None,         random_state=None, ...),         n_jobs=-1,         param_grid={'learning_rate': [0.01, 0.1, 0.25],         'max_depth': [5, 10, 15],         'n_estimators': [10, 35, 100]},         scoring='accuracy')</pre>	<pre>In [116]: xgb_gs.best_params_  Out[116]: {'learning_rate': 0.1, 'max_depth': 10, 'n_estimators': 100}</pre> <p>Checking model fitness</p> <hr/> <p>Train score: 0.9329 Test score: 0.8616</p>

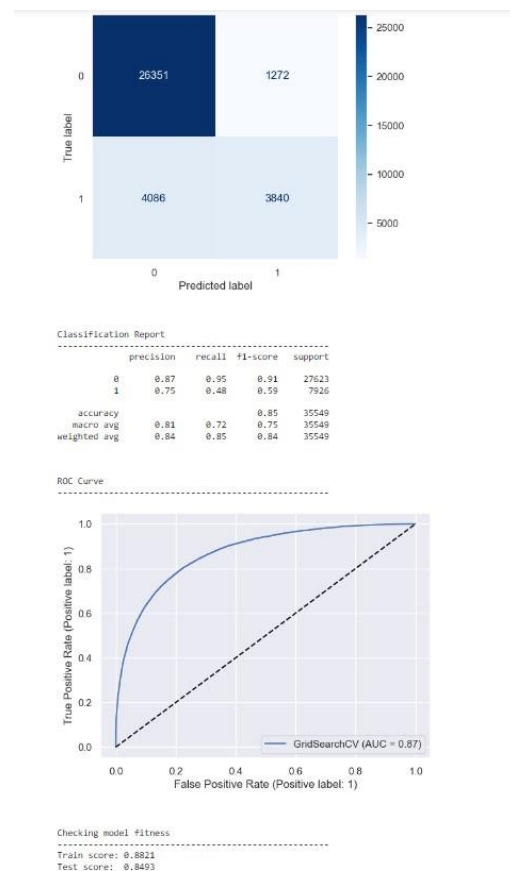
## Performance Metrics Comparison Report (2 Marks):

Model	Baseline Metric	Optimized Metric
Logistic Regression	 <pre> Classification Report -----               precision    recall  f1-score   support       0       0.92      0.88      0.86      27623      1       0.52      0.77      0.62       7926   accuracy      0.79      35549  macro avg     0.72      0.78      0.74      35549  weighted avg   0.83      0.79      0.80      35549           </pre>  <pre> ROC Curve -----           </pre> <pre> Checking model fitness ----- Train score: 0.7886 Test score: 0.7914           </pre> <p><b>Accuracy: 79%    AUC: 0.86</b></p>	 <pre> Classification Report -----               precision    recall  f1-score   support       0       0.87      0.94      0.90      27623      1       0.70      0.50      0.59       7926   accuracy      0.79      35549  macro avg     0.79      0.72      0.74      35549  weighted avg   0.83      0.84      0.83      35549           </pre>  <pre> ROC Curve -----           </pre> <pre> Checking model fitness ----- Train score: 0.8409 Test score: 0.8421           </pre> <p><b>Accuracy: 84%    AUC: 0.86</b></p>

## Random Forest




Accuracy: 85%    AUC: 0.88



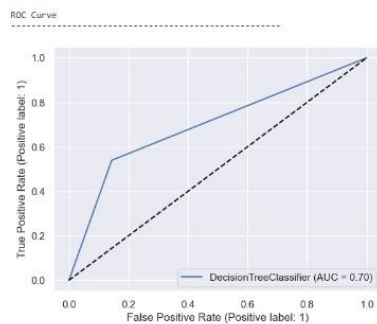
Accuracy: 85%    AUC: 0.87

## Decision Tree



Classification Report


	precision	recall	f1-score	support
0	0.87	0.86	0.86	27623
1	0.52	0.54	0.53	7926
accuracy				0.79
macro avg	0.69	0.70	0.69	35549
weighted avg	0.79	0.79	0.79	35549



Checking model fitness


Train score: 1.0  
Test score: 0.7855

Accuracy: 79%    AUC: 0.70



Classification Report

	precision	recall	f1-score	support
0	0.86	0.95	0.90	27623
1	0.73	0.45	0.56	7926
accuracy				0.84
macro avg	0.79	0.70	0.73	35549
weighted avg	0.83	0.84	0.83	35549

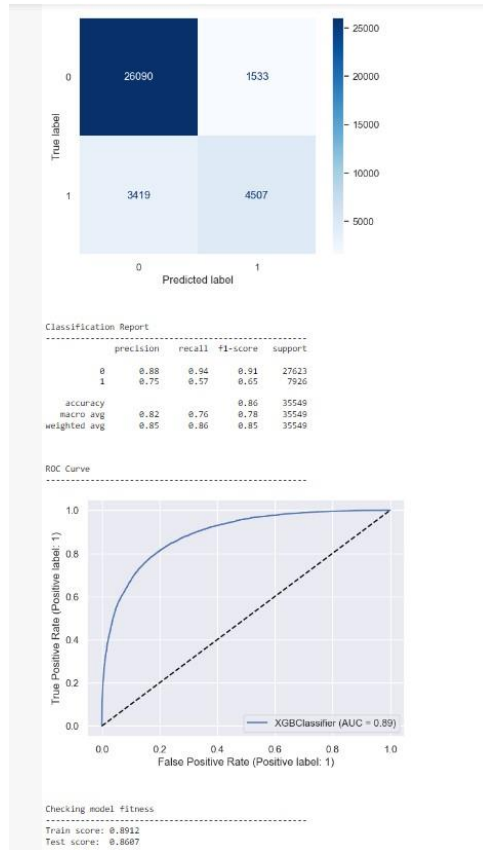


Checking model fitness

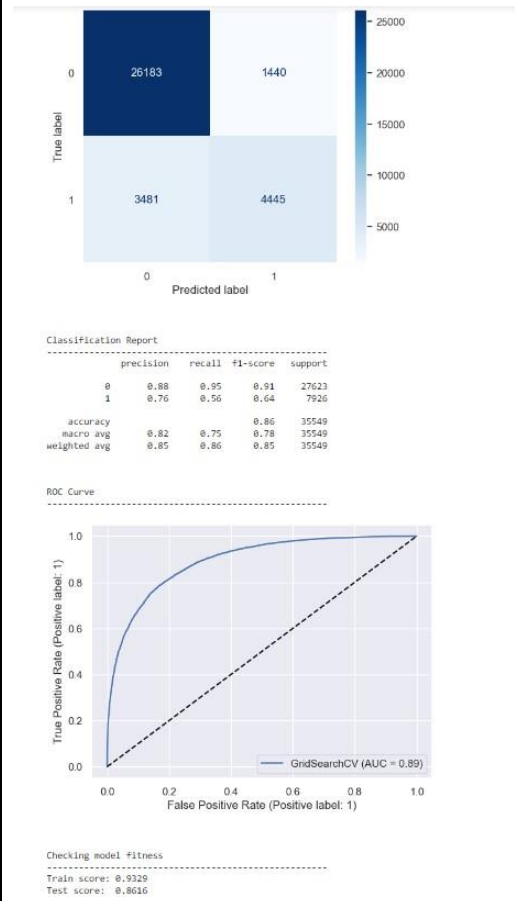
Train score: 0.8475  
Test score: 0.8485

Accuracy: 84%    AUC: 0.85

## XG-Boost



Accuracy: 86%    AUC: 0.89



Accuracy: 86%    AUC: 0.89

**Final Model Selection Justification (2 Marks):**

Final Model	Reasoning
XG-Boost	The best performing model is the hyperparameter-tuned XG-Boost model with an accuracy of approximately 86%. The scores for both the training and testing data were similar, reducing concerns of the model being overfit.